

Workshop No.2 Systems Sciences

Juan Santiago Ramos Ome, Arlo Nicolas Ocampo

May 14, 2025

1 Introduction

Develop an autonomous agent capable of learning and adapting to a simulated environment using reinforcement learning. The agent will monitor two intersections with traffic lights with multiple traffic participants (cars, motorcycles, and pedestrians).

An autonomous agent will be designed to control and manage the traffic lights in two intersections, with the goal of optimizing traffic flow with reinforcement learning in the simulation, the agent learns to respond to the traffic lights by observing the traffic conditions, in order to minimize waiting time and avoid traffic into the two points.

This workshop delves into the analysis and refinement of an autonomous agent designed for traffic light control, utilizing concepts from dynamical systems. The primary goal is to enhance the agent's ability to adapt to complex traffic scenarios by employing tools such as Markov Decision Processes (MDP) and reinforcement learning, specifically Q-learning. Furthermore, the workshop explores the use of phase portraits to visualize system behavior and emphasizes the importance of stability and convergence in the agent's learning process. The design of advanced feedback mechanisms is also addressed to enable the agent to respond effectively to uncertain and dynamic environments.

2 Update timeline

The following timeline presents an updated version of the previously submitted work plan. This revision incorporates refinements based on project progress, received feedback, and a clearer definition of development stages. While maintaining a weekly structure, the updated timeline introduces more detailed technical tasks and specific objectives for each phase, particularly in the implementation of the simulation environment, reinforcement learning agent training, and full system integration.

Week 1 – Introduction to Systems and Problem Analysis

- Introduction to System Theory, interactions, flows, synergy, emergent properties
- Overview of holistic approaches and system thinking
- Introduction to complex systems behavior and system engineering
- Scenario analysis: two intersections with traffic lights
- Identification of key actors (vehicles, sensors, pedestrians, environment)
- **Objective:** Understand and document the problem using a system thinking approach

Week 2 – Theoretical Foundations of AI and Control Systems

- Introduction to AI, intelligent systems, machine learning
- Explain main goals of AI, symbolic vs subsymbolic AI
- Introduction to feedback mechanisms, feedback loops, open vs closed loop systems
- Concept of homeostasis, dynamic systems, and chaos theory
- Review of supervised, unsupervised, and reinforcement learning
- Literature review of related AI applications in traffic management
- **Objective:** Connect system theory and AI principles to the project context

Week 3 – Environment and Gym Design

- Define environment states (vehicles, people, time)
- Define actions (traffic light changes)
- Logical structure of environment, rewards, sensors
- Begin implementation of the Gym environment
- **Objective:** Create the functional simulation space using Python and Gymnasium

Week 4 – Sensor Logic and Traffic Simulation

- Vehicle simulation programming
- Traffic light logic and timers
- Sensor modeling (zone counting, observation space)
- **Objective:** Establish simulation and observation infrastructure

Week 5 – Q-Learning Setup and Initial Training

- Implement basic RL agent with Q-table
- Training in multiple simple scenarios
- Visualize agent behavior and state transitions
- **Objective:** Initial training with Q-learning and basic performance review

Week 6 – Evaluate and Tune Q-Learning

- Analyze Q-learning performance
- Refine reward function
- Tune parameters and visualize learning improvements
- **Objective:** Improve the Q-learning agent through testing and feedback

Week 7 – Transition to Deep Q-Network (DQN)

- Introduce Stable-Baselines3
- Implement DQN agent
- Train and compare performance with Q-table agent
- **Objective:** Deploy the first DQN-based model

Week 8 – DQN Performance Analysis

- Analyze DQN results
- Adjust rewards, learning rates, and experience replay settings
- **Objective:** Stabilize and improve the DQN agent

Week 9 – Deep Dive into Agent Optimization

- Optimize hyperparameters
- Review impact of changes on traffic efficiency
- **Objective:** Refine agent decision-making capabilities

Week 10 – Multi-Agent and Edge Case Consideration

- Consider expansion to multi-agent settings
- Handle edge cases (pedestrian surge, sudden vehicle increase)
- **Objective:** Prepare system for more complex, realistic scenarios

Week 11 – Compare Learning Models

- Evaluate and compare Q-learning and DQN quantitatively
- Metrics: average wait time, traffic flow, agent decisions
- **Objective:** Provide evidence-based performance comparison

Week 12 – Incorporate Additional Complexity

- Add complexity: emergency vehicles, weather, time of day
- Modify agent behavior and reward structures
- **Objective:** Increase simulation realism and agent robustness

Week 13 – System Integration

- Integrate all modules: sensors, simulation, learning, control logic
- Begin system-wide testing
- **Objective:** Full system integration

Week 14 – Final Testing and Analysis

- Run simulations and record data
- Evaluate system performance and edge-case handling
- Final tweaks and preparations for documentation
- **Objective:** Validate system under all planned test conditions

Week 15 – Documentation and Delivery

- Final report (including system diagrams and code documentation)
- Presentation preparation
- Submit deliverables
- **Objective:** Deliver complete project with results and analysis

3 System Dynamics Analysis:

3.1 Mathematical/Simulation Model

For the agent to learn to intelligently control traffic lights, it is necessary to represent mathematically the environment and its decisions. Two principal models that we are used for this: the Markov Decision Process (MDP) and the Q-learning algorithm.

The MDP allows the operation of the traffic system to be described as a series of different states (how many vehicles there are, the state of the traffic light, and actions like changing or maintaining the traffic light phase, and rewards like reducing the waiting time). This helps to define what the agent observes and what it can do.

Then the Q-learning is applied, a reinforcement learning algorithm that allows the agent to learn through trial and error. With the simulation the agent tests different actions, observes the results, and learns which decisions are best for improving vehicular flow. Over time, the agent learns a strategy that optimizes traffic light behavior.

MDP

Is a decision mathematical framework for describing decision-making problems in dynamic and uncertain environments, such as traffic.

$$MDP = (S, A, P, R) \quad (1)$$

Where:

- S: Describe the current situation of the environment, in these case, the traffic and traffic lights.
- A: They are the decisions that the agent can make in a given state.
- P: Describes the probability of the system moving from one state to another.
- R: It is a numerical value that represents how good an action was in a given state.

2.2 Phase Portraits or Diagrams

Phase portraits are a way to visualize the behavior of a dynamical system in its state space. The "state space" is the set of all possible states that the system can be in. In our traffic control scenario, a "state" could be defined by variables like:

- Number of vehicles before traffic light A
- Number of vehicles between traffic lights A and B
- Number of vehicles after traffic light B
- Average waiting time at each light
- Current phase of each traffic light (red, yellow, green)

And a probable phase portrait for this project could be a portrait where:

- X-axis = Number of vehicles before traffic light A
- Y-axis = Average waiting time at traffic light B

Attractors would represent desirable traffic flow patterns. For example:

- A stable attractor might be a point with "moderate vehicle numbers before A" and "low average waiting time at B." This indicates efficient traffic flow.
- The agent's goal is to guide the system towards these attractor states by making appropriate traffic light decisions.

In traffic control, chaos could manifest as:

- Unexpected surges in traffic that lead to oscillations in waiting times and vehicle numbers.

Swift changes could be present as:

- A long arrow might represent a sudden increase in waiting time due to a traffic incident.
- The agent's actions (e.g., changing traffic light phases) should ideally produce smooth transitions in the phase portrait, avoiding abrupt changes that could disrupt traffic flow.

Phase portraits can show how the system responds to different inputs or conditions.

- Possible scenarios: "high congestion", "low demand", "different arrival rates of vehicles," etc.

- Comparing these phase portraits will reveal how the agent adapts (or fails to adapt) to varying inputs.

Stability is indicated by trajectories that remain within a bounded region of the phase portrait. Convergence is when trajectories approach an attractor over time.

- The agent should promote stability, preventing traffic conditions from spiraling out of control (e.g., unbounded queues).
- Convergence means that the agent learns to consistently guide the system towards optimal traffic flow.

4 Feedback Loop Refinement:

4.1 Enhanced Control Mechanisms

Table 1: Additional sensors : Following the feedback from the previous workshop, these are the sensors that will help collect the best data on the road.

sensors	What does the sensor do?	Data provided by the sensor
Speed sensor	Measures the average speed of vehicles in each lane	Average speed (km/h or m/s)
Vehicle distance sensor	Estimate traffic congestion in real time	Average distance between vehicles
People waiting sensor	How long a person has been waiting	People waiting time
Number of cars passing per second	Measures the number of cars that pass the traffic light	Average number of cars passing per second

Table 2: More granular rewards : The reward functions that proposed in the previous workshop were highly complex. The new functions that will allow the agent to earn rewards and improve its efficiency will be shown below.

Reward signal	Type of score reward	Advantage
Fluency Reward	+1 . crossing vehicles	the number of vehicles crossing increases
Penalty for prolonged waiting of people	-1 . waiting people	Prevents the system from ignoring the people
You have to cross cars when it's green	-2 points if there is a green light but no vehicle is crossing	Improves efficient use of green time

These rewards will allow the agent to function well, allowing a constant flow of vehicles to avoid congestion and allowing the people to cross the street. In the first reward, we see that there is a reward for each vehicle that crosses the traffic light. However, a time limit will be set for waiting for the people (if there are any). If this limit is exceeded, points will begin to be deducted. If there are no people, the flow of vehicles will continue.

4.2 Stability and Convergence

Stability and convergence are critical concepts for our traffic control agent. We need to ensure that the agent not only performs well but also does so reliably and consistently.

- **Stability Criteria:**

- **Bounded Behavior:** The system should exhibit bounded behavior, meaning that key metrics like vehicle queue lengths and waiting times remain within acceptable limits. The agent should prevent the system from reaching extreme states like gridlock.
- **Resilience:** The agent should be resilient to disturbances, such as sudden changes in traffic flow or sensor noise, and maintain stable performance.

- **Convergence:**

- **Performance Improvement:** Convergence refers to the agent's ability to improve its performance over time. In our case, this means that through learning, the agent should progressively reduce congestion and waiting times.
- **Steady State:** Ideally, the agent's performance should converge to a steady state, where further learning does not significantly change its effectiveness.

- **Evaluation:**

- **Theoretical Evaluation:** We will use the underlying mathematical framework (Markov Decision Process) and reinforcement learning principles (Q-learning, DQN) to provide a theoretical basis for stability and convergence.
- **Practical Evaluation:**
 - * **Simulation Analysis:** We will analyze simulation results to measure stability and convergence. This involves tracking metrics like average waiting time, vehicle throughput, and queue lengths over extended simulation periods.
 - * **Scenario Testing:** Testing the agent under various scenarios (high/low traffic, incidents) will further validate its stability.

5 Iterative Design Outline:

4.1 Update to Project Plan

To effectively incorporate dynamic systems concepts and achieve more advanced agent behavior, the project plan needs to be updated with the following:

- **Enhanced Data Structures:**
 - **Time Series Data:** We need to handle time series data to capture the temporal nature of traffic flow. This includes storing sequences of traffic conditions, agent actions, and performance metrics.
 - **Dynamic State Representation:** Instead of simple static state representations, we may need more complex structures (e.g., graphs) to represent the dynamic relationships between different parts of the traffic network.
- **Advanced Algorithms:**
 - **Recurrent Neural Networks (RNNs):** To capture temporal dependencies in traffic patterns, we can explore using RNNs or LSTM networks.
 - **Kalman Filtering:** Implement Kalman filters to estimate traffic flow parameters and predict future states, enhancing the agent's ability to adapt to uncertainty.
 - **Exploration-Exploitation Strategies:** Implement more sophisticated exploration strategies in reinforcement learning to balance exploring new actions and exploiting known good ones.
- **Frameworks and Tools:**
 - **Deep Learning Frameworks:** Integrate deep learning frameworks like TensorFlow or PyTorch to support more complex models like RNNs and DQN variants.
 - **Traffic Simulation Tools:** Consider using more advanced traffic simulation tools (e.g., SUMO) to create highly realistic and complex traffic scenarios.
- **Additional Planning:**
 - **Computational Resources:** Account for the increased computational demands of more complex models and simulations.
 - **Validation:** Emphasize rigorous validation techniques to ensure the agent's robustness and generalization ability.
 - **Feedback Loop Design:** Refine the design of feedback loops to enable the agent to adapt more rapidly to changing conditions.

Table 4: Simulation parameters

Parameter	Description	Example
Simulation duration	Total time of training	30 min to 1 hour
Decision making	How many seconds the agent act	Every 30 seconds
times of change of states	Minimum time to stay green/red light	30 seconds

Scenario variations : To test the adaptability of the program, it will be create scenarios such as:

- High congestion, to see if the agent invites traffic jams
- Low demand for vehicles, to see if the agent optimizes the use of the traffic light
- Weather conditions, see if the agent adapts phases for safety

References

- [1] Gymnasium Documentation. [Online]. Available: <https://gymnasium.farama.org/index.html>
- [2] Stable-Baselines3 Documentation – Reliable Reinforcement Learning Implementations. [Online]. Available: <https://stable-baselines3.readthedocs.io/en/master/>
- [3] R. Lin, “Create your own environment using the OpenAI Gym library — GridWorld,” *Medium*, Mar. 3, 2022. [Online]. Available: <https://reneelin2019.medium.com/create-your-own-environment-using-the-openai-gym-library-gridworld-8ca9f18e00a4>
- [4] Vitality Learning, “Solving the Taxi Problem Using OpenAI Gym and Reinforcement Learning,” *Medium*, Nov. 15, 2024. [Online]. Available: <https://vitalitylearning.medium.com/solving-the-taxi-problem-using-openai-gym-and-reinforcement-learning-0317e089b48f>
- [5] “Tutorial: An Introduction to Reinforcement Learning Using Open AI Gym.” [Online]. Available: <https://www.gocoder.one/blog/rl-tutorial-with-openai-gym/>
- [6] Gymnasium Documentation. [Online]. Available: <https://gymnasium.farama.org/index.html>
- [7] R. Lin, “Create your own environment using the OpenAI Gym library GridWorld,” *Medium*, [Online]. Available: <https://reneelin2019.medium.com/create-your-own-environment-using-the-openai-gym-library-gridworld-8ca9f18e00a4>

- [8] M. Newman, *Networks*. Oxford University Press, 2018.
- [9] S. H. Strogatz, *Nonlinear Dynamics and Chaos: With Applications to Physics, Biology, Chemistry, and Engineering*. CRC Press, 2018.
- [10] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*. MIT Press, 2018.